


```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from itertools import product
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
```


```
df = pd.read_csv('/content/heart.csv')
df
```



	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2
...	...	...	...	...	...	...	...	...	...	...	...	...	...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2


303 rows × 14 columns

```
df.info()
```




```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
df.shape
```




```
(303, 14)
```

```
df.columns
```



```
Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
       'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
      dtype='object')
```

```
null_values = df.isnull().sum()
print("Null values in each column:\n", null_values)
```



```
Null values in each column:
age      0
sex      0
cp       0
trestbps 0
```

```
chol      0
fbs       0
restecg   0
thalach   0
exang     0
oldpeak   0
slope     0
ca        0
thal      0
target    0
dtype: int64
```


```
X = df.drop("target", axis=1)
y = df["target"].apply(lambda x: 1 if x > 0 else 0) # Binarize the target
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Hyperparameters: n\_estimators, max\_depth, min\_samples\_split, and max\_leaf\_nodes.

```
# Define hyperparameters and their values
hyperparameters = {
    'n_estimators': [100, 500],
    'max_depth': [10, 20],
    'min_samples_split': [2, 10],
    'max_leaf_nodes': [10, 20]
}
```

```
# Hyperparameter Optimazation using GridSearch
model = RandomForestClassifier()
model_gs = GridSearchCV(estimator=model, param_grid=hyperparameters)
model_gs.fit(X_train, y_train)
```



```
# Get the best parameters
model_gs.best_params_
```

```
{'max_depth': 20,
 'max_leaf_nodes': 10,
 'min_samples_split': 2,
 'n_estimators': 100}
```

```
# Get the best score
model_gs.best_score_
```

```
0.8304421768707482
```

## Grid Search

```
# Random Forest with Grid Search
y_pred_gs = model_gs.predict(X_test)

print("\nRandom Forest Grid Search Performance:")
print("Accuracy:", accuracy_score(y_pred_gs, y_test))
print("\nClassification Report:")
print(classification_report(y_pred_gs, y_test))
```

```
Random Forest Grid Search Performance:
Accuracy: 0.8852459016393442

Classification Report:
              precision    recall  f1-score   support

0               0.86         0.89         0.88         28
1               0.91         0.88         0.89         33
```

accuracy			0.89	61
macro avg	0.88	0.89	0.88	61
weighted avg	0.89	0.89	0.89	61


## Random Search

```
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV
```

```
# Initialize the RandomForestClassifier
rf = RandomForestClassifier()
```

```
# Initialize RandomizedSearchCV
random_search = RandomizedSearchCV(
    estimator=rf,
    param_distributions=hyperparameters,
    n_iter=100, # Number of parameter settings that are sampled
    cv=5, # 5-fold cross-validation
    verbose=2,
    random_state=42,
    n_jobs=-1 # Use all available cores
)
```

```
# Fit RandomizedSearchCV to the data
random_search.fit(X_train, y_train)
```

 Fitting 5 folds for each of 16 candidates, totalling 80 fits  
 /usr/local/lib/python3.10/dist-packages/sklearn/model\_selection/\_search.py:305: UserWarning:

```
warnings.warn(
> RandomizedSearchCV
> estimator: RandomForestClassifier
> RandomForestClassifier
```

```
from sklearn.metrics import precision_score, recall_score, f1_score, confusion_matrix, accuracy_score
```

```
# Print the best parameters and the corresponding score
print(f"Best parameters found: {random_search.best_params_}")
print(f"Best cross-validation score: {random_search.best_score_}")
```


```
# Predict with the best estimator
best_rf = random_search.best_estimator_
y_pred = best_rf.predict(X_test)
```

```
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Test set accuracy: {accuracy}")
```

```
# Calculate and print precision, recall, f1 score, and support
precision = precision_score(y_test, y_pred, average=None)
recall = recall_score(y_test, y_pred, average=None)
f1 = f1_score(y_test, y_pred, average=None)
cm = confusion_matrix(y_test, y_pred)
```

```
# Calculate support from confusion matrix
support = cm.sum(axis=1)
```

```
# Print the classification report manually
print("\nClassification Report:")
print(f"{'Class':<10}{'Precision':<10}{'Recall':<10}{'F1-Score':<10}{'Support':<10}")
for i in range(len(precision)):
    print(f"{'i':<10}{precision[i]:<10.2f}{recall[i]:<10.2f}{f1[i]:<10.2f}{support[i]:<10}")
```

 Best parameters found: {'n\_estimators': 100, 'min\_samples\_split': 10, 'max\_leaf\_nodes': 10, 'max\_depth': 10}  
 Best cross-validation score: 0.8263605442176871  
 Test set accuracy: 0.8852459016393442

```
Classification Report:
Class      Precision Recall    F1-Score  Support
```

0	0.89	0.86	0.88	29
1	0.88	0.91	0.89	32